

Disk Schedulers for Solid State Drives

Jaeho Kim
School of Computer
Science University of
Seoul Seoul, Korea
kjhnet07@uos.ac.kr

Yongseok Oh
School of Computer
Science University of
Seoul Seoul, Korea
ysoh@uos.ac.kr

Eunsam Kim
School of Computer &
Inform. Eng. Hongik
University Seoul, Korea
eskim@hongik.ac.kr

Jongmoo Choi
Division of Information
and CS Dankook
University Seoul, Korea
choijm@dankook.ac.kr

Donghee Lee*
School of Computer
Science University of
Seoul Seoul, Korea
dhl_express@uos.ac.kr

Sam H. Noh
School of Computer &
Inform. Eng. Hongik
University Seoul, Korea
samhnoh@hongik.ac.kr

Outline

- Introduction
- Flash Translation Layer
- Linux Disk I/O Scheduler
- Logical Block
- Design of New Disk Schedulers
- Experimental Results
- Conclusions

Introduction

- The SSD (Solid State Drive) has been introduced and is gaining popularity in embedded systems and laptops
- Because of the differences in device characteristics, the current schedulers may not adequately schedule requests for SSDs

Flash Translation Layer (1/3)

- Solve the constraint that page overwrite is forbidden
- Maintain a pool of writable pages by pre-erasing the redundant blocks
- Approach
 - Page mapping
 - Block mapping

Flash Translation Layer (2/3)

- Page mapping
 - The map translates a sector number to a combination of page number and block number where the sector data exists
 - Be excellent random write performance
 - Increase garbage collection overhead after randomly write

Flash Translation Layer (3/3)

- Block mapping
 - The map translates a sector number to a combination of page number and block number where the sector data exists
 - To modify data, the block mapping FTL reserves some redundant blocks called **log blocks**


Linux Disk I/O Scheduler

- Noop
 - Merge adjacent requests to a larger one
- Deadline
 - If the waiting time of a request exceeds its deadline, the request is served immediately
- Anticipatory
 - Waits for new in-coming read requests for a predetermined period
- CFQ (Complete Fair Queuing)
 - Separate queue for each process and serves requests of queues in round-robin order


Logical Block (1/4)

- Read performance of an SSD is almost consistent and independent of the ordering and geometrical distance between read requests
- Sequential writes to a logical block is much more efficient than random writes going to various logical blocks

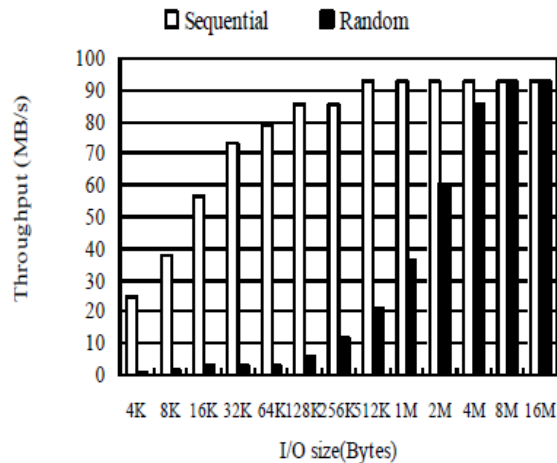
Logical Block (2/4)




(a) Samsung SSD read



(b) Mtron SSD read



(c) Samsung SSD write




(d) Mtron SSD write


Logical Block (3/4)

- LBA-bundle (Logical Block Aligned-bundle)
 - A way to maximize write performance in an SSD
 - Arrange write requests into bundles the size of a logical block so that write requests falling in a logical block belong to the same bundle
- The writing order of the bundles themselves will not be important if all requests within the bundle are written sequentially at a time

Logical Block (4/4)



When the merger occur



Design of New Disk Schedulers (1/2)

- The IRBW-FIFO (Individual Read Bundled Write FIFO) Scheduler
 - Apply FIFO ordering to the bundles of write requests and individual read requests.

Design of New Disk Schedulers (2/2)

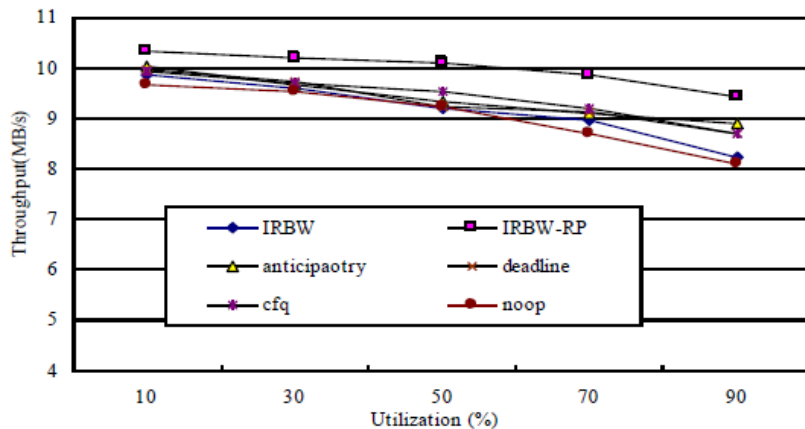
- The IRBW-FIFO-RP (Individual Read Bundled Write FIFO with Read Preference) Scheduler
 - Maintains separate FIFO ordering among read requests and among the bundles of write requests
 - Gives higher priority to read requests than to the bundles of write requests
 - To avoid write starvation, allow each LBA-bundle to yield to a read request only once

Experimental Results (1/3)

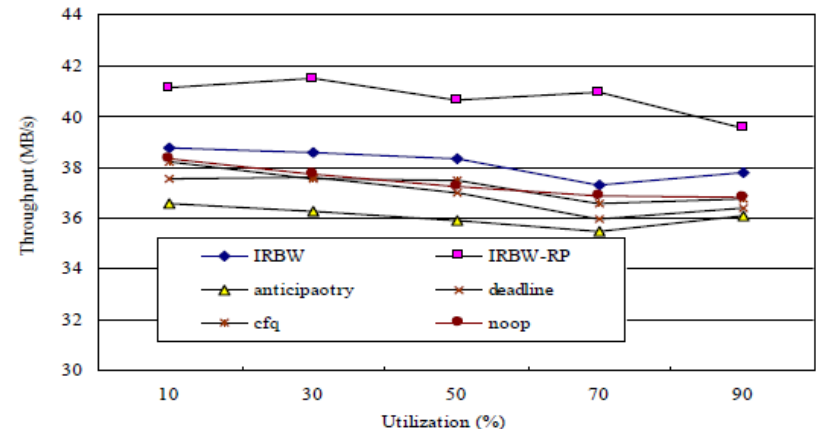
- Throughput results
- Response time results
- Experimental environment

Type	Specifics	
CPU/RAM	Intel Core2 Duo E4500 2.2GHz / 2GB DRAM	
SSD	Samsung 64GB MCCOE64G5MPP SATA-2	
	Mtron 16GB MSD-SATA6025 SATA-1	
OS	Linux-kernel 2.6.23 / Ext3 File system	
Benchmark	Postmark	Type-A: 1KB~32KB file size
		Type-B: 1MB~16MB file size
	IOmeter	File server access pattern
Targets	IRBW-FIFO, IRBW-FIFO-RP, and existing Linux I/O schedulers	

Experimental Results (2/3) - Throughput Results




(a) Postmark Type-A

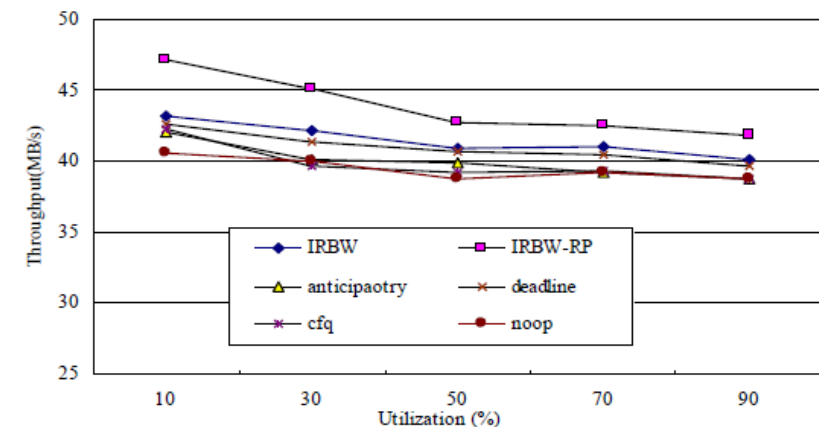


(b) Postmark Type-B

Figure 8. Throughput of Postmark benchmarks on Samsung SSD



(a) Postmark Type-A




(b) Postmark Type-B

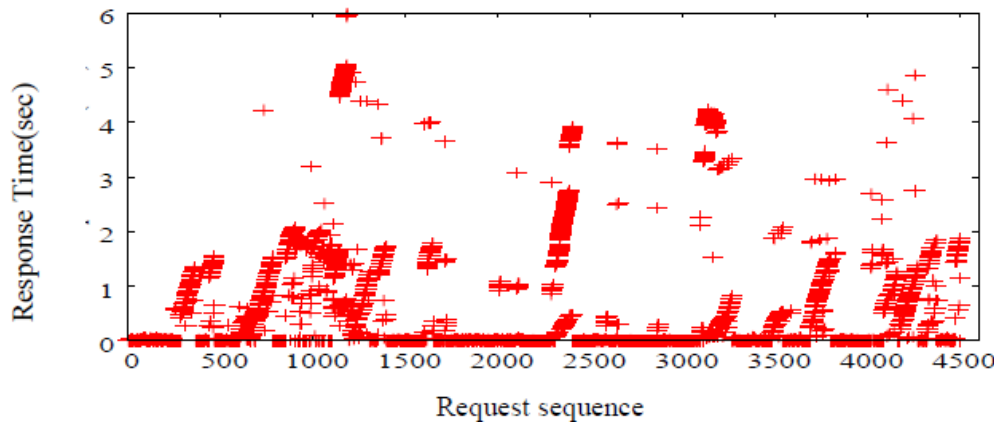
Figure 9. Throughput of Postmark benchmarks on Mtron SSD

Experimental Results (3/3) - Response Time Results

- Average response times of the two schedulers are 0.475 and 0.458



(a) Deadline I/O scheduler



(b) IRBW-FIFO-RP I/O scheduler

Figure 11. Response time of Postmark benchmark on Samsung SSD

Conclusions

- SSDs have much faster read service times than the magnetic disks with the service times being almost constant, while write request service times are more complex
- IRBW-FIFO and IRBW-FIFO-RP arrange write requests into LBA-bundles while reads are independently scheduled